How to build an homogeneous language for heterogeneous platforms

You don't have to trade abstraction for control

Julien Richard-Foy   Olivier Barais   Jean-Marc Jézéquel

IRISA, Université de Rennes 1

{first}.{last}@irisa.fr

Abstract

Writing large Web applications is known to be difficult. One challenge comes from the fact that the application's logic is scattered into heterogeneous clients and servers, making it difficult to share code between both sides or to move code from one side to the other. Another challenge is performance: while Web applications rely on ever more code on the client-side, they may run on smart phones with little hardware capabilities. These two challenges raise the following problem: how to benefit from high-level languages and libraries making code complexity easier to manage and abstracting over the clients and servers differences without trading this engineering comfort for performance? This article presents high-level abstractions defined as deep embedded DSLs in Scala, that can (1) generate efficient code leveraging the target platform characteristics, (2) be shared between client and server code. Code written with our DSL has a high level of abstraction but runs as fast as hand tuned low-level JavaScript code.

Categories and Subject Descriptors   D.3.3 [*Programming Languages*]: Language Constructs and Features

General Terms   Languages, Software Engineering

Keywords   Heterogeneous code generation, Domain-specific languages

1. Introduction

Web applications are attractive because they require no installation or deployment steps on clients and enable large scale collaborative experiences. However, writing large Web applications is known to be difficult [14, 16]. One challenge comes from the fact that the business logic is scattered into heterogeneous client-side and server-side environments [12, 17]. This gives less flexibility in the engineering process and requires a higher maintenance effort: if a piece of logic is implemented on client-side and finally needs to be implemented on server-side instead, the code can not be reused and the feature needs to be completely rewritten (and *vice versa*). Even worse, logic parts that run on both client-side and server-side need to be duplicated. For instance, HTML fragments may be built from the server-side when a page is requested by a client, but they may also be built from the client-side to perform an incremental update subsequent to an user action. How could developers write HTML fragment definitions once and render them on both client-side and server-side?

The more interactive the application is, the more logic needs to be duplicated between the server-side and the client-side, and the higher is the complexity of the client-side code. Developers use libraries and frameworks to get high-level abstractions on client-side, making their code easier to reason about and to maintain, but also making their code run less efficiently (due to *abstraction penalty*).

Performance is a primary concern in Web applications, because they are expected to run on a broad range of devices, from the powerful desktop personal computer to the less powerful smart phone. "Every 100 ms delay costs 1% of sales", said Amazon in 2006.

Using the same programming language on both server-side and client-side could improve the software engineering process by enabling code reuse between both sides. Incidentally, the JavaScript language – which is currently the most supported the action language on Web clients – can be used on server-side, and an increasing number of programming languages or compiler back-ends can generate JavaScript code (*e.g.*Java/GWT [4], SharpKit[1], Dart [7], Kotlin[2], ClojureScript [13], Fay[3], Haxe [3] or Opa[4]).

However, using the same programming language is not enough because the client and server programming environments are not the same. For instance, DOM fragments can be defined on client-side using the standard DOM API, but this API does not exist on server-side. How to define a common vocabulary for such concepts? And how to make the executable code leverage the native APIs, when possible, for performance reasons?

Generating efficient code for heterogeneous platforms is hard to achieve in an extensible way: the translation of common abstractions like collections into their native counterpart (JavaScript arrays on client-side and standard library's collections on server-side) may be hard coded in the compiler, but that would not scale to handle all the abstractions a complete application may use (*e.g.*HTML fragment definitions, form validation rules, or even some business data type that may be represented differently).

On one hand, for engineering reasons, developers want to write Web applications using a single high-level language, abstracting over the target platforms differences and reducing code complexity.

---

[1] http://sharpkit.net

[2] http://kotlin.jetbrains.org/

[3] http://fay-lang.org/
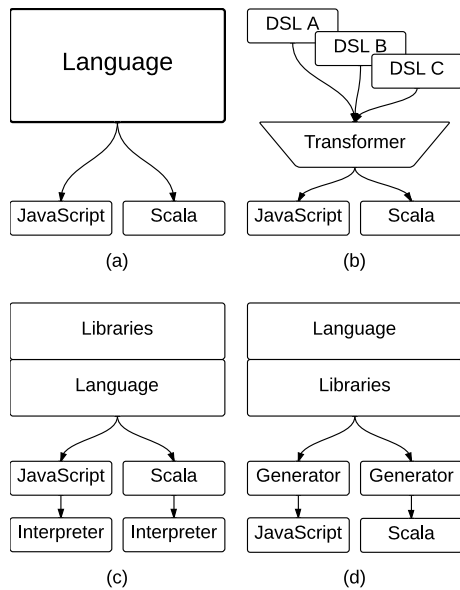
[4] http://opalang.org/

Figure 1. Language engineering processes

But on the other hand, for performance reasons, they want to keep control on the way their code is compiled to each target platform. How to solve this dilemma?

Compiled domain-specific embedded languages [6] allow the definition of domain-specific languages (DSLs) as libraries on top of a host language, and to compile them to a target platform. The deep embedding gives the opportunity to control the code generation scheme for a given abstraction and target platform.

Kossakowski *et al.* introduced *js-scala*, a compiled embedded DSL defined in Scala that generates JavaScript code, making it possible to write the client-side code of Web applications using JavaScript [11]. However, Kossakowski *et al.* did not address the engineering dilemma described above. This paper enriches *js-scala* with the following contributions:

- we define high-level abstractions typically used in Web programming generating low-level code leveraging native Web browsers APIs;

- in the case of abstractions shared between servers and clients, we specialize the code generation in order to leverage the target platform native environments.

We validate our approach with a case study implemented with various candidate technologies and discuss the relative pro and cons of them. Though the code written in our DSL is high-level and can be shared between clients and servers, it has the same runtime performances as hand-tune low-level JavaScript code.

The remainder of this paper is organized as follows. The next section introduces existing approaches for defining cross-compiling languages. Sections 3 and 4 present our contribution. Section 5 gives implementation details. Section 6 evaluates the contribution. Section 7 concludes.

2. Background

This section presents different approaches for defining cross-platform programming languages.

Fat Languages The first approach for defining a cross-platform language consists in hard-coding, in the compiler, the code gen-

eration scheme of each language feature to each target platform. Figure 1 (a) depicts this process. In order to support a feature related to a specific domain, the whole compiler pipeline (parser, code generator, *etc.*) may have to be adapted. This approach gives *fat* languages because a lot of concepts are defined at the language level: general programming concepts such as naming, functions, classes, and more specific concepts such as HTML fragment definition. Examples of such languages are Links [5], Opa and Dart [7]. These languages are difficult to extend because each concept is defined in the compiler, and modifying a compiler requires a high effort. Furthermore, these languages also require to support common programming abstraction and composition mechanisms, as general purpose languages do. So they usually try to re-invent the features of general purpose languages. We argue that this approach for defining programming languages is difficult to scale: for every problem you have to rewrite a full-featured programming language besides addressing the concepts specific to the problem domain.

Domain-Specific Languages Another approach consists in defining several independent domain-specific languages [19], each one focusing on concerns specific to a given problem domain, and then to combine all the source artifacts written with these language into one executable program, as shown in figure 1 (b). Defining such a language requires a minimal effort compared to the previous approach. On the other hand, it is difficult to have interoperability between DSLs (ref). [8] shows such a domain-specific language for defining Web applications.

Thin Languages Alternatively, one can define concepts relative to a specific domain as libraries on top of a thin general purpose language (this is also referred to as *embedded* domain specific languages). Figure 1 (c) depicts this approach. The general purpose language is used as a host language and does not need to be modified if a new concept is introduced, because concepts are defined as pure libraries. However, this approach gives no opportunity to translate a concept efficiently according to the target platform characteristics. Examples of languages following this approach are Java/GWT, Kotlin, HaXe and SharpKit.

Deeply Embedded Languages The last approach, shown in figure 1 (d), can be seen as a middle-ground between the two previous approaches: DSLs are embedded in a host language but use a code generation process. (benefits and limitations of this approach)

Lightweight Modular Staging [18] is a framework for defining deeply embedded DSLs in Scala using staging [10]. It has been used to define high-performance DSLs for parallel computing [2].

3. High-Level Abstraction for Client-Side Code: Selectors

In a Web application, the user interface is defined by a HTML document that can be updated by the JavaScript code. A typical operation consists in searching some "interesting" element in the document, in order to extract its content, replace it or listen to user events triggered on it (such as a mouse click). The standard API provides several functions to search elements in a HTML document according to their name or attribute values. Figure 2 summarizes the available functions and their differences.

The `querySelector` and `querySelectorAll` are the most general functions while the others handle special cases. For the developer it is not convenient to have to master several functions performing similar tasks. In fact, most JavaScript developers use the jQuery library [1][5] that provides only one high-level func-

---

[5] According to the following document: http://trends.builtwith.com/javascript, jQuery is used by more than 40% of the top million sites

Function	Description
`querySelector(s)`	First element matching the CSS selector `s`
`getElementById(i)`	Element which attribute `id` equals to `i`
`querySelectorAll(s)`	All elements matching the CSS selector `s`
`getElementsByTagName(n)`	All elements of type `n`
`getElementsByClassName(c)`	All elements which `class` attribute contains `c`

Figure 2. Standard selectors API

Listing 1. Selectors in plain JavaScript
```
function getWords() {
  var form = document.getElementById('add-user');
  var sections =
    form.getElementsByTagName('fieldset');
  var results = [];
  for (var i = 0 ; i < sections.length ; i++) {
    var words = sections[i]
      .getElementsByClassName('word');
    results[i] = words;
  }
  return results
}
```

Listing 2. Selectors in jQuery
```
function getWords() {
  var form = $('#add-user');
  var sections = $('fieldset', form);
  return sections.map(function () {
    return $('.word', this)
  })
}
```

Listing 3. Selectors in js-scala
```
def getWords() = {
  val form = document.find("#add-user")
  val sections = form.findAll("fieldset")
  sections map (_.findAll(".word"))
}
```

tion to search for elements. Listings 1 and 2 show two equivalent JavaScript programs performing element searches, the first one using the native APIs and the second one using jQuery.

jQuery provides an API that is simpler to master because it has less functions, but by doing so it can not benefit from the performance of the browser's implementation of specialized search functions (`getElementById`, `getElementsByTagName` and `getElementsByClassName`).

Listing 3 shows how to implement listing 2 using *js-scala*. We provide two functions for searching elements: `find` to find the first element matching a selector and `findAll` to find all the matching elements. During staging these functions try to analyze the selector that is passed as parameter and, when appropriate, produce code using the specialized API, otherwise they produce code using `querySelector` and `querySelectorAll`. As a result, listing 3 generates a JavaScript program identical to listing 1: the high-level abstraction (the `find` and `findAll` functions) exist only in the initial program, not in the final JavaScript program. It gives two advantages: (1) the execution of the final program is more performant

Listing 4. Unsafe code
```
var loginWidget =
  document.querySelector("div.login");
var loginButton =
  loginWidget.querySelector("button.submit");
loginButton.addEventListener("click", handler);
```

Listing 5. Defensive programming to handle null references
```
var loginWidget =
  document.querySelector("div.login");
if (loginWidget !== null) {
  var loginButton =
    loginWidget.querySelector("button.submit");
  if (loginButton !== null) {
    loginButton.
      addEventListener("click", handler);
  }
}
```

because of the use of specialized APIs, (2) the final program's size is smaller because it does not need to include jQuery. (Move this last sentence in the discussion)

4. High-Level Abstractions Shared on Clients and Servers

4.1 Monads Sequencing

Null references are a known source of problems in programming languages [9, 15]. For example, consider listing 4 finding a particular widget in the page and then a particular button in the widget. The native `querySelector` method returns `null` if no node matched the given selector in the document. If we run this code in a page where the widget is not present, it will throw an error and stop further JavaScript execution. Defensive code can handle `null` reference, but leads to very cumbersome code, as shown in listing 5.

Some programming languages encode optional values with a monad (*e.g.* `Maybe` in Haskell and `Option` in Scala). In that case, sequencing over the monad encodes optional value dereferencing. If the language supports a convenient syntax for monad sequencing, it brings a convenient syntax for optional value dereferencing, alleviating developers from the burden of defensive programming.

Listing 6 implements in *js-scala* a program equivalent to listing 5. The `for` notation is used to sequence computations over optional values encoded with a monad.

We want to define a DSL that has both the safety and performance of listing 5 but the expressiveness of listing 4. We can get safety by wrapping potentially null values of type `Rep[A]` in a container of type `Rep[Option[A]]` requiring explicit dereferencing, we can get expressiveness by using the Scala `for` notation for dereferencing, and finally we can get performance by generating code that does not actually wraps values in a container but instead checks if they are `null` or not when dereferenced. The wrapping container exists only at the stage-level and is removed during the code generation. Listing 6 shows how to write the same program using our DSL (implementation details are given in section 5). The evaluation of this listing produces a graph of statements from which JavaScript code equivalent to listing 5 is generated.

4.2 DOM Fragments Definition

In this section we show how we can define a template engine as an embedded DSL with minimal effort. This template engine is statically typed and able to insert dynamic content in a safe way. It provides a powerful expression language, requires no extra

Listing 6. Handling null references in js-scala

```
for {
  loginWidget <- document.find("div.login")
  loginButton <- loginWidget.find("submit.button")
} loginButton.on(Click)(handler)
```

Listing 7. JavaScript DOM API

```
var articleUi = function (article) {
  var div = document.createElement('div');
  div.setAttribute('class', 'article');
  var span = document.createElement('span');
  var name =
    document.createTextNode(article.name + ': ');
  span.appendChild(name);
  div.appendChild(span);
  var strong = document.createElement('strong');
  var price = document.createTextNode(article.price);
  strong.appendChild(price);
  div.appendChild(strong);
  return div
};
```

Listing 8. HTML

```
<div class=article>
  <span>French wine: </span>
  <strong>10</strong>
</div>
```

Listing 9. DOM definition DSL

```
def articleUi(article: Rep[Article]) =
  el('div, 'class -> 'article)(
    el('span)(article.name + ": "),
    el('strong)(article.price)
  )
```

compilation step and can be used on both client-side and server-side.

Because the template engine is defined as en embedded DSL, we can reuse Scala's constructs:

- a function taking some parameters and returning a DOM fragment directly models a template taking parameters and returning a DOM fragment ;

- the type system type-checks template definitions and template calls ;

- the Scala language itself is the expression language ;

- compiling a template is the same as compiling user code.

So the only remaining work consists in defining the DSL vocabulary to define DOM nodes. We provide a `el` function to define a tag and any `String` value is considered to be a text node.

Listing 9 uses our DSL and generates a code equivalent to listing 7. The readability has been highly improved: nesting tags is just like nesting code blocks, HTML entities are automatically escaped in text nodes, developers have the full computational power of Scala to inject dynamic data and DOM fragments definitions are written using functions so they compose just as functions compose. These benefits come with no performance loss because the DSL generates code building DOM fragments by using the native JavaScript API.

Our DSL is equivalent to a template engine with Scala as the expression language. Making it usable on both server and client sides was surprisingly as simple as defining another code generator for the DSL, producing Scala code.

Listing 10. Selectors optimization

```
def find(receiver: Exp[Selector],
         selector: Exp[String]) =
  getConstIdentifier(selector) match {
    case Some(id) if receiver == document =>
      DocumentGetElementById(Const(id))
    case _ =>
      SelectorFind(receiver, selector)
  }
```

Listing 11. Null reference handling DSL code generator

```
trait JSGenOptionOps extends JSGenEffect {
  val IR: OptionOpsExp
  import IR._
  override def emitNode(sym: Sym[Any], rhs: Def[Any]) =
    rhs match {
      case OptionIsEmpty(o) =>
        emitValDef(sym, q" $o === null")
      case OptionForeach(o, a, b) =>
        stream.println(q"if ($o !== null) {")
        emitValDef(a, quote(o))
        emitBlock(b)
        stream.println("}")
      case _ =>
        super.emitNode(sym, rhs)
    }
}
```

For instance, the template written in listing 9 produces the following Scala code usable on server-side:

```
def articleUi(article: Article) = {
  val x0 = <span>{ article.name + ": " }</span>
  val x1 = <strong>{ article.price }</strong>
  val x2 =
    <div class="article">
      {List(x0, x1)}
    </div>
  x2
}
```

We are able to tackle the code sharing issues described in the introduction (more details?) because of the embdedded nature of our DSLs: dynamic content of templates is written using embedded DSLs too, so their translation into JavaScript and Scala is managed by their respective code generators.

5. Implementation
[6]

5.1 Selectors

Listing 10

5.2 Null references

Code generation consists in traversing the statement nodes produced by the program evaluation according to their dependencies and to emit the code corresponding to each statement. LMS already sorts the statements graph so DSL authors just need to say how to emit code for each statement node of their DSL. Listing 11 shows such a code generator for the `null` reference handling DSL. The `emitNode` method handles `OptionIsEmpty` and `OptionForeach` nodes. In the case of the `OptionIsEmpty` node, it simply generates an expression testing if the value is `null`, in the case of the `OptionForeach` node, it wraps the code block dereferencing the value within a `if` checking that the value is not `null`.

---

[6] The code is available at http://github.com/js-scala

Figure 3. Benchmarks on a real application

5.3 DOM Fragments

6. Evaluation

We implemented several applications using js-scala. We also have written several implementations of a complete application, using different approaches to write the client and server sides, and compared the amount of code written, the runtime performances and the ability to modularize the code and to maintain it.

6.1 Real World Application

Chooze [7] is an existing complete application for making polls. It allows users to create a poll, define the choice alternatives, share the poll, vote and look at the results. It contains JavaScript code to handle the dynamic behavior of the application: double-posting prevention, dynamic form update and rich interaction with the document.

The application was initially written using jQuery. We rewrote it using several technologies for the client-side part: plain JavaScript (without third-party library), js-scala, GWT and HaXe.

Vanilla JavaScript Low-level code.

jQuery High-level code.

js-scala High-level code. HTML fragment definition reused between server and client sides.

HaXe Low-level code.

GWT High-level code?

6.2 Benchmarks, Code Metrics

Our goal is to evaluate the level of abstraction provided by each solution and their performances. We took the number of lines of code as a measure of the level of abstraction. We also measured the size of the data sent to the client-side. We also measured the ability to share code between client and server sides.

Rather than doing a micro-benchmark focusing on just one abstraction, we performed a global benchmark that is more likely to reflect the real performances of the application: we simulated user actions on a Web page (2000 clicks on buttons, triggering a dynamic update of the page and involving the use of the Option monad, the selectors API and the HTML fragment definition API) and measured the time it took complete them. We only measured the execution time of client-side code execution. The tests were run on a DELL Latitude E6430 laptop with 8 GB of RAM, on the Google Chrome v27 Web browser. Figure 3 shows the benchmark results.

7. Conclusion

We implemented a high-level language abstracting over client and server heterogeneity but producing efficient code.

---

[7] http://chooze.herokuapp.com

References

[1] B. Bibeault and Y. Kats. *jQuery in Action*. Dreamtech Press, 2008.

[2] K. J. Brown, A. K. Sujeeth, H. J. Lee, T. Rompf, H. Chafi, M. Odersky, and K. Olukotun. A heterogeneous parallel framework for domain-specific languages. In *Parallel Architectures and Compilation Techniques (PACT), 2011 International Conference on*, pages 89–100. IEEE, 2011.

[3] N. Cannasse. Using haxe. *The Essential Guide to Open Source Flash Development*, pages 227–244, 2008.

[4] P. Chaganti. *Google Web Toolkit: GWT Java Ajax Programming*. Packt Pub Limited, 2007.

[5] E. Cooper, S. Lindley, P. Wadler, and J. Yallop. Links: Web programming without tiers. In *Formal Methods for Components and Objects*, pages 266–296. Springer, 2007.

[6] C. Elliott, S. Finne, and O. De Moor. Compiling embedded languages. *Journal of Functional Programming*, 13(3):455–481, 2003.

[7] R. Griffith. The dart programming language for non-programmers-overview. 2011.

[8] D. Groenewegen, Z. Hemel, L. Kats, and E. Visser. Webdsl: A domain-specific language for dynamic web applications. In *Companion to the 23rd ACM SIGPLAN conference on Object-oriented programming systems languages and applications*, pages 779–780. ACM, 2008.

[9] T. Hoare. Null references: The billion dollar mistake. *Presentation at QCon London*, 2009.

[10] U. Jørring and W. L. Scherlis. Compilers and staging transformations. In *Proceedings of the 13th ACM SIGACT-SIGPLAN symposium on Principles of programming languages*, pages 86–96. ACM, 1986.

[11] G. Kossakowski, N. Amin, T. Rompf, and M. Odersky. JavaScript as an Embedded DSL. In J. Noble, editor, *ECOOP 2012 – Object-Oriented Programming*, volume 7313 of *Lecture Notes in Computer Science*, pages 409–434, Berlin, Heidelberg, 2012. Springer Berlin Heidelberg. doi: 10.1007/978-3-642-31057-7_19. URL https://github.com/js-scala/js-scala/.

[12] J. Kuuskeri and T. Mikkonen. Partitioning web applications between the server and the client. In *Proceedings of the 2009 ACM symposium on Applied Computing*, SAC '09, pages 647–652, New York, NY, USA, 2009. ACM. ISBN 978-1-60558-166-8. doi: 10.1145/1529282.1529416. URL http://doi.acm.org/10.1145/1529282.1529416.

[13] M. McGranaghan. Clojurescript: Functional programming for javascript platforms. *Internet Computing, IEEE*, 15(6):97–102, 2011.

[14] T. Mikkonen and A. Taivalsaari. Web applications - spaghetti code for the 21st century. In *Proceedings of the 2008 Sixth International Conference on Software Engineering Research, Management and Applications*, pages 319–328, Washington, DC, USA, 2008. IEEE Computer Society. ISBN 978-0-7695-3302-5. doi: 10.1109/SERA.2008.16. URL http://dl.acm.org/citation.cfm?id=1443226.1444030.

[15] M. Nanda and S. Sinha. Accurate interprocedural null-dereference analysis for java. In *Software Engineering, 2009. ICSE 2009. IEEE 31st International Conference on*, pages 133–143. IEEE, 2009.

[16] J. C. Preciado, M. L. Trigueros, F. Sánchez-Figueroa, and S. Comai. Necessity of methodologies to model rich internet applications. In *WSE*, pages 7–13. IEEE Computer Society, 2005. ISBN 0-7695-2470-2.

[17] R. Rodríguez-Echeverría. Ria: more than a nice face. In *Proceedings of the Doctolral Consortium of the International Conference on Web Engineering*, volume 484. CEUR-WS.org, 2009.

[18] T. Rompf. *Lightweight Modular Staging and Embedded Compilers*. PhD thesis, IC, Lausanne, 2012. URL http://library.epfl.ch/theses/?nr=5456.

[19] A. Van Deursen, P. Klint, and J. Visser. Domain-specific languages: an annotated bibliography. *ACM Sigplan Notices*, 35(6):26–36, 2000.